


libdispatch - event handling


- Grand Central Dispatch
- Asynchronous & concurrent programming model
- From apple
- <http://libdispatch.macosforge.org/>


chenj@lemote.com


What is event handling?


- Example: event handling in glib
 - Create a GMainContext
 - Attach one or more GSources
 - GSource: wrap event and callback
 - GSource:pollfd → 1:n
 - Built-in GSource:
 - timeout source
 - child watch source
 - idle source
 - ...
 - Create a GMainLoop associating with GMainContext, then "g_main_loop_run"














Travel through pending_dispatches:

For each source:

1. Remove **ready flag**
2. If isn't a **can_recurse** source, block it
3. Mark as **in_call**
4. Link to a list which lives in stack:
GMainDispatch->dispatching_sources
5. Run dispatch, destroy the source later if returns false
6. Cleanup:
 - 1> Unlink from *GMainDispatch->dispatching_sources*
 - 2> Restore the **in_call** state
 - 3> Unblock if is blocked

prepare()

dispatch()

Dispatching

Initial[n]

Event handling in glib

- Call *g_main_context_iteration* or *g_main_loop_run in callback* in callback:
 - **can_recurse** GSource
- Running callback will slow down event handling – Long running callback leads to bad responsiveness

Event handling in libdispatch

- `dispatch_source_t`: wrap event and callback
 - source:kevent → n:1
 - Set a event callback and (optional) cancel callback
 - Set source's target queue – callback runs in target queue
 - Source's priority is determined by its target queue
- Monitor event and do early process in *mgr queue*
 - In a single thread of the highest priority thread pool
 - Based on kevent
 - No priority in monitoring and early processing stage

Event handling: libdispatch vs glib

- libdispatch has better responsiveness
 - Running callback will not slow down event handling
 - kevent is more efficient than poll
- glib can run multi-GMainContexts in multi-threads

Create/Setup source


- `dispatch_source_create`
 - `dispatch_source_t` inherits from `dispatch_queue_t`
 - Specify source's type – `DISPATCH_SOURCE_TYPE_DATA_ADD` , `DISPATCH_SOURCE_TYPE_DATA_OR` , `DISPATCH_SOURCE_TYPE_MACH_RECV` , `DISPATCH_SOURCE_TYPE_MACH_SEND` , `DISPATCH_SOURCE_TYPE_PROC` , `DISPATCH_SOURCE_TYPE_READ` , `DISPATCH_SOURCE_TYPE_SIGNAL` , `DISPATCH_SOURCE_TYPE_TIMER` , `DISPATCH_SOURCE_TYPE_VNODE` , `DISPATCH_SOURCE_TYPE_WRITE`
 - Created in SUSPEND state
 - `dispatch_source_set_event_handler_f`

Run source(Installation stage1)

- `dispatch_resume`
 - `_dispatch_wakeup`
 - send source to its target queue
- Invoke source in target queue (`_dispatch_queue_invoke`)
 - `_dispatch_source_invoke` → redirect to *mgr queue*
- Wake up *mgr queue* (`_dispatch_wakeup`)


Wake up mgr queue

`_dispatch_wakeup(&_dispatch_mgr_q)`




Wake up mgr queue

`_dispatch_wakeup(&_dispatch_mgr_q)`




Wake up mgr queue

`_dispatch_wakeup(&_dispatch_mgr_q)`




Wake up mgr queue

`_dispatch_wakeup(&_dispatch_mgr_q)`




`_dispatch_update_kq(const struct kevent *ke)`




Wake up mgr queue

`_dispatch_wakeup(&_dispatch_mgr_q)`




`_dispatch_update_kq(const struct kevent *ke)`




Invoke mgr queue

`_dispatch_get_kq_init`


- Create kqueue
- Set monitoring EVFILT_USER
- Send mgr queue to its target queue ||  `_dispatch_queue_invoke`

Invoke mgr queue

`_dispatch_cqueue_invoke`



Event loop – running timers



- Process timer sources
- Two types of timer: wall timer; abs timer
- Link timers of the same type together(by order of expiration time, early → late)
- For each timer list:
 - For each expired timer:
 - Early process
 - Update next expiration time
 - Adjust its position in list
 - Wake up it

Event loop – Calc timeout



- Final timeout is 0 if exist an expired timer.
- Or calc by travelling through two timer lists:
 1. Find the earliest expiration times of all valid timers(not SUSPEND and has set expiration time) for each timer list
 2. Calc timeout from expiration time
 3. Convert abs timeout to wall timeout, find an smaller one

Event loop – kevent

`_dispatch_mgr_invoke`



`_dispatch_run_timers`

`_dispatch_get_next_timer_fire`

`/* select workaround stuff */`

`kevent(_dispatch_kq, NULL, 0,
kev, 1, timeoutp);`

`_dispatch_mgr_thread2`

- A generic method of kernel event notification on BSD.
- Usage:
 - `int kq = kqueue()`
 - `int ret = kevent(int kq,
struct kevent *changelist,
int nchanges,
struct kevent *eventlist,
int nevents,
const struct timespec *timeout);`

struct kevent


- **struct kevent**

<code>uintptr_t</code>	ident;	Value used to identify this event. The exact interpretation is determined by the attached filter, but often is a file descriptor.
<code>int16_t</code>	filter;	Identifies the kernel filter used to process this event.
<code>uint16_t</code>	flags;	Actions to perform on the event.
<code>uint32_t</code>	fflags;	Filter-specific flags.
<code>intptr_t</code>	data;	Filter-specific data value.
<code>void</code>	*udata;	Opaque user-defined value passed through the kernel unchanged.


filters

- EVFILT_READ
Data available for read on monitored fd
- EVFILT_WRITE
Buffer available for write on monitored fd
- EVFILT_VNODE
Notify operations on inode of monitored fd
- EVFILT_PROC
- EVFILT_SIGNAL
- EVFILT_TIMER
- ...


Event loop – process event




Event loop – process event




Event loop – process event



Run source(Installation stage2)



Run source(Installation stage2)



Run source(Installation stage2)

_dispatch_queue_drain

Run source(Installation stage2)

`_dispatch_queue_drain`



`_dispatch_queue_invoke((dispatch_queue_t) source)`



Run source(Installation stage2)

`_dispatch_queue_drain`



`_dispatch_queue_invoke((dispatch_queue_t) source)`



`_dispatch_kevent_merge`

- `ds->ds_is_installed = true`
- Try to merge monitored kevents (each kevent can be identified by filter and ident)
- if needs update, call:

`_dispatch_kevent_resume`

Run source(Installation stage2)

`_dispatch_queue_drain`



`_dispatch_queue_invoke((dispatch_queue_t) source)`



`dx_invoke`

`_dispatch_source_invoke`



`_dispatch_kevent_merge`

- `ds->ds_is_installed = true`
- Try to merge monitored kevents (each kevent can be identified by filter and ident)
- if needs update, call:


`_dispatch_kevent_resume`





`kevent(...)`

`_dispatch_update_kq`



Event loop – process event





Event loop – process event



Event loop – process event



Event loop – process event



- Early process
- Wake up source

Process event(early)

- Classify events:
 - **level**: `ds_pending_data = kev->data`
 - `DISPATCH_SOURCE_TYPE_READ`
 - `DISPATCH_SOURCE_TYPE_WRITE`
 - **adder**: `ds_pending_data += ke->data`
 - `DISPATCH_SOURCE_TYPE_SIGNAL`
 - `DISPATCH_SOURCE_TYPE_TIMER`
 - **or**: `ds_pending_data |= (kev->fflags & ds->ds_pending_data_mask)`
 - `DISPATCH_SOURCE_TYPE_VNODE`

Process event

- Wake up source → send to its target queue
- `_dispatch_queue_drain`
↓
`_dispatch_queue_invoke((dispatch_queue_t) source)`



`dx_invoke`

`_dispatch_source_invoke`



`_dispatch_source_latch_and_call`

- `ds_handler_func(ds_handler_ctxt);`

END